

Progetto

Andrea Corvaglia

17 agosto 2019

TO DO:

La Gpu con 112 livelli è inutilizzabile, possiamo dividere in integrata e dedicata, oppure dividere per le marche.

@Dario: Per la gpu farei innanzitutto una divisione tra integrata e dedicata, e tra quelle dedicate le suddividerei per VRAM (inizialmente pensavo di cercare dei prezzi per avere una stima quantitativa di quanto valga la gpu dedicata, però bisognerebbe cercare prezzi medi e visto che i modelli son tanti mi pare troppo lavoro e non chissà quanto significativo)

Anche la memoria è problematica, ho già estratto la presenza o meno dell'SSD, ma bisogna pensare ad un modo di valutare la memoria.

@Dario: ottimo la divisione SSD/HD, per avere un indicazione quantitativa potremmo cercare il prezzo di un gb di SSD e di un gb di HD da una fonte affidabile, e pesare l'importanza delle dimensioni su questo indicatore, che ne dite?

```
data <- read.csv("../data/Laptop2.csv")
str(data)
```

```
## 'data.frame':    1303 obs. of  17 variables:
## $ X              : int  1 2 3 4 5 6 7 8 9 10 ...
## $ Company        : Factor w/ 19 levels "Acer","Apple",...: 2 2 8 2 2 1 2 2 3 1 ...
## $ Product        : Factor w/ 618 levels "110-15ACL (A6-7310/4GB/500GB/W10)",...: 302 300 51 302 302 ...
## $ TypeName       : Factor w/ 6 levels "2 in 1 Convertible",...: 5 5 4 5 5 4 5 5 5 5 ...
## $ Inches         : num  13.3 13.3 15.6 15.4 13.3 15.6 15.4 13.3 14 14 ...
## $ ScreenResolution: Factor w/ 40 levels "1366x768","1440x900",...: 24 2 9 26 24 1 26 2 9 16 ...
## $ Cpu            : Factor w/ 118 levels "AMD A10-Series 9600P 2.4GHz",...: 55 53 64 75 57 15 74 53 ...
## $ Ram           : int   8 8 8 16 8 4 16 8 16 8 ...
## $ Memory        : Factor w/ 39 levels "1.0TB HDD","1.0TB Hybrid",...: 5 3 17 30 17 27 16 16 30 17 ...
## $ Gpu           : Factor w/ 110 levels "AMD FirePro W4190M",...: 59 52 54 10 60 18 61 52 98 62 ...
## $ OpSys         : Factor w/ 9 levels "Android","Chrome OS",...: 5 5 6 5 5 7 4 5 7 7 ...
## $ Weight        : num   1.37 1.34 1.86 1.83 1.37 2.1 2.04 1.34 1.3 1.6 ...
## $ Price         : num   1340 899 575 2537 1804 ...
## $ Frequenza     : num    2.3 1.8 2.5 2.7 3.1 3 2.2 1.8 1.8 1.6 ...
## $ Risoluzione   : Factor w/ 15 levels "1366x768","1440x900",...: 11 2 4 13 11 1 13 2 4 4 ...
## $ Pixel         : int  4096000 1296000 2073600 5184000 4096000 1049088 5184000 1296000 2073600 2073600 ...
## $ SolidStateDisk : Factor w/ 2 levels "False","True": 2 1 2 2 2 1 1 1 2 2 ...
```

```
head(data)
```

```
##   X Company      Product TypeName Inches
## 1 1  Apple MacBook Pro Ultrabook  13.3
## 2 2  Apple Macbook Air Ultrabook  13.3
## 3 3    HP      250 G6  Notebook  15.6
## 4 4  Apple MacBook Pro Ultrabook  15.4
## 5 5  Apple MacBook Pro Ultrabook  13.3
## 6 6   Acer   Aspire 3  Notebook  15.6
##                                ScreenResolution      Cpu Ram
```

```
## 1 IPS Panel Retina Display 2560x1600 Intel Core i5 2.3GHz 8
## 2 1440x900 Intel Core i5 1.8GHz 8
## 3 Full HD 1920x1080 Intel Core i5 7200U 2.5GHz 8
## 4 IPS Panel Retina Display 2880x1800 Intel Core i7 2.7GHz 16
## 5 IPS Panel Retina Display 2560x1600 Intel Core i5 3.1GHz 8
## 6 1366x768 AMD A9-Series 9420 3GHz 4
## Memory Gpu OpSys Weight
## 1 128GB SSD Intel Iris Plus Graphics 640 macOS 1.37
## 2 128GB Flash Storage Intel HD Graphics 6000 macOS 1.34
## 3 256GB SSD Intel HD Graphics 620 No OS 1.86
## 4 512GB SSD AMD Radeon Pro 455 macOS 1.83
## 5 256GB SSD Intel Iris Plus Graphics 650 macOS 1.37
## 6 500GB HDD AMD Radeon R5 Windows 10 2.10
## Price Frequenza Risoluzione Pixel SolidStateDisk
## 1 1339.69 2.3 2560x1600 4096000 True
## 2 898.94 1.8 1440x900 1296000 False
## 3 575.00 2.5 1920x1080 2073600 True
## 4 2537.45 2.7 2880x1800 5184000 True
## 5 1803.60 3.1 2560x1600 4096000 True
## 6 400.00 3.0 1366x768 1049088 False
```

```
summary(data)
```

```
## X Company Product
## Min. : 1.0 Dell :297 XPS 13 : 30
## 1st Qu.: 331.5 Lenovo :297 Inspiron 3567 : 29
## Median : 659.0 HP :274 250 G6 : 21
## Mean : 660.2 Asus :158 Legion Y520-15IKBN: 19
## 3rd Qu.: 990.5 Acer :103 Vostro 3568 : 19
## Max. :1320.0 MSI : 54 Inspiron 5570 : 18
## (Other):120 (Other) :1167
## TypeName Inches
## 2 in 1 Convertible:121 Min. :10.10
## Gaming :205 1st Qu.:14.00
## Netbook : 25 Median :15.60
## Notebook :727 Mean :15.02
## Ultrabook :196 3rd Qu.:15.60
## Workstation : 29 Max. :18.40
##
## ScreenResolution
## Full HD 1920x1080 :507
## 1366x768 :281
## IPS Panel Full HD 1920x1080 :230
## IPS Panel Full HD / Touchscreen 1920x1080: 53
## Full HD / Touchscreen 1920x1080 : 47
## 1600x900 : 23
## (Other) :162
## Cpu Ram
## Intel Core i5 7200U 2.5GHz :190 Min. : 2.000
## Intel Core i7 7700HQ 2.8GHz:146 1st Qu.: 4.000
## Intel Core i7 7500U 2.7GHz :134 Median : 8.000
## Intel Core i7 8550U 1.8GHz : 73 Mean : 8.382
## Intel Core i5 8250U 1.6GHz : 72 3rd Qu.: 8.000
## Intel Core i5 6200U 2.3GHz : 68 Max. :64.000
## (Other) :620
```

```
##          Memory          Gpu
## 256GB SSD      :412  Intel HD Graphics 620 :281
## 1TB HDD        :223  Intel HD Graphics 520 :185
## 500GB HDD      :132  Intel UHD Graphics 620 : 68
## 512GB SSD      :118  Nvidia GeForce GTX 1050: 66
## 128GB SSD + 1TB HDD: 94  Nvidia GeForce GTX 1060: 48
## 128GB SSD      : 76  Nvidia GeForce 940MX  : 43
## (Other)       :248  (Other)              :612
##      OpSys      Weight      Price      Frequenza
## Windows 10:1072  Min.    :0.690  Min.    : 174  Min.    :0.900
## No OS      : 66  1st Qu.:1.500  1st Qu.: 599  1st Qu.:2.000
## Linux      : 62  Median :2.040  Median : 977  Median :2.500
## Windows 7  : 45  Mean   :2.039  Mean   :1124  Mean   :2.299
## Chrome OS  : 27  3rd Qu.:2.300  3rd Qu.:1488  3rd Qu.:2.700
## macOS     : 13  Max.   :4.700  Max.   :6099  Max.   :3.600
## (Other)    : 18
##      Risoluzione      Pixel      SolidStateDisk
## 1920x1080:841  Min.    :1049088  False:460
## 1366x768 :308  1st Qu.:1440000  True :843
## 3840x2160: 43  Median :2073600
## 3200x1800: 27  Mean   :2168807
## 1600x900 : 23  3rd Qu.:2073600
## 2560x1440: 23  Max.   :8294400
## (Other)      : 38
```

```
nums <- sapply(data, is.numeric)
var_numeric <- data[,nums]
head(var_numeric)
```

```
##      X Inches Ram Weight      Price Frequenza      Pixel
## 1 1    13.3   8    1.37 1339.69      2.3 4096000
## 2 2    13.3   8    1.34  898.94      1.8 1296000
## 3 3    15.6   8    1.86  575.00      2.5 2073600
## 4 4    15.4  16    1.83 2537.45      2.7 5184000
## 5 5    13.3   8    1.37 1803.60      3.1 4096000
## 6 6    15.6   4    2.10  400.00      3.0 1049088
```


```
data$Weight<-as.numeric(data$Weight)
data$Ram<-as.numeric(data$Ram)
```

```
sapply(data, function(x)(sum(is.na(x))))
```

```
##          X          Company      Product      TypeName
##          0              0              0              0
##      Inches ScreenResolution      Cpu      Ram
##          0              0              0              0
##      Memory          Gpu      OpSys      Weight
##          0              0              0              0
##      Price      Frequenza      Risoluzione      Pixel
##          0              0              0              0
## SolidStateDisk
##          0
```

```
# Non ci sono missing data!
```

```
plot(data$Company,data$Price)
```




```
class(data$Ram)
```

```
## [1] "numeric"
```

```
plot(density(data$Frequenza))
```

density.default(x = data\$Frequenza)




```
#hist(data$Price, breaks=25, probability=TRUE)
#lines(density(data$Price))
```

```
library(ggplot2)

## Registered S3 methods overwritten by 'ggplot2':
##   method      from
## [.quosures    rlang
## c.quosures     rlang
## print.quosures rlang

ggplot(data,aes(x = Price)) +
  geom_histogram(aes(y = ..density..),
    bins= 25,
    fill = "grey",
    color ="black") +
  geom_vline(xintercept = quantile(data$Price, 0.50), color = "dark red", lty = 2) +
  geom_vline(xintercept = mean(data$Price), color = "dark blue", lty = 2) +
  labs(x = "Price", y ="Density") +
  ggtitle("Price Distribution with mean and median") +
  geom_density()
```

Price Distribution with mean and median




Quite skewed to the right, mean > media

We could try to apply a correction like $\log(Y)$

```
data$LogPrice=log(data$Price)
ggplot(data,aes(x = log(Price))) +
  geom_histogram(aes(y = ..density..),
    bins= 25,
    fill = "grey",
    color ="black") +
```


```
geom_vline(xintercept = quantile(data$LogPrice, 0.50), color = "dark red", lty = 2) +
geom_vline(xintercept = mean(data$LogPrice), color = "dark blue", lty = 2) +
labs(x = "log(Price)", y = "Density") +
ggtitle("log(Price) Distribution with mean and median")+ geom_density()
```



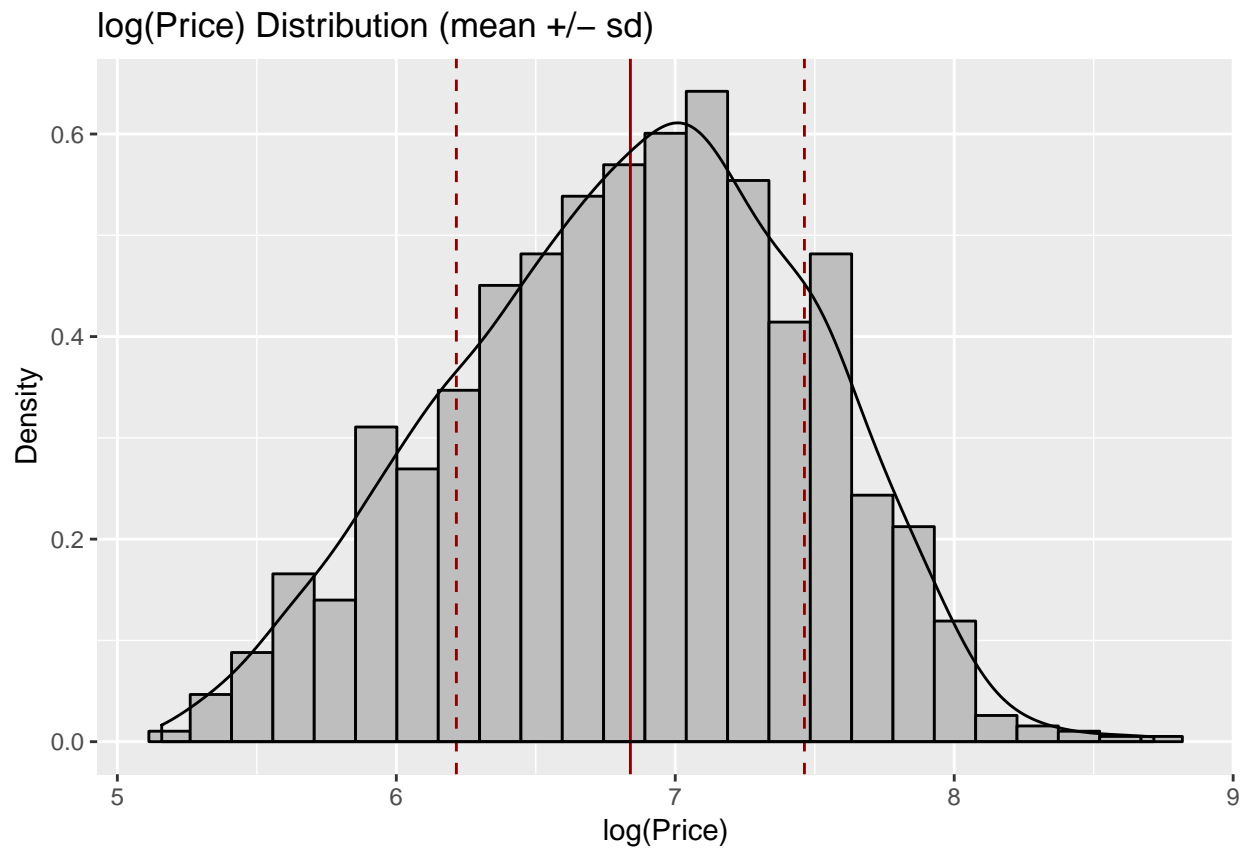
Now the distribution is looking a bit better (as regards normality)

```
ggplot(data,aes(x = Price)) +
  geom_histogram(aes(y =..density..),
    bins= 25,
    fill = "grey",
    color ="black") +
  geom_vline(xintercept = mean(data$Price), color = "dark red") +
  geom_vline(xintercept = mean(data$Price) + sd(data$Price), color = "dark red", lty = 2) +
  geom_vline(xintercept = mean(data$Price) - sd(data$Price), color = "dark red", lty = 2) +
  labs(x = "Price", y ="Density") +
  ggtitle("Price Distribution (mean +/- sd)") +
  geom_density()
```

Price Distribution (mean \pm sd)




```
ggplot(data,aes(x = log(Price))) +
  geom_histogram(aes(y =..density..),
    bins= 25,
    fill = "grey",
    color ="black") +
  geom_vline(xintercept = mean(data$LogPrice), color = "dark red") +
  geom_vline(xintercept = mean(data$LogPrice) + sd(data$LogPrice), color = "dark red", lty = 2) +
  geom_vline(xintercept = mean(data$LogPrice) - sd(data$LogPrice), color = "dark red", lty = 2) +
  labs(x = "log(Price)", y ="Density") +
  ggtitle("log(Price) Distribution (mean +/- sd)") +
  geom_density()
```




```
ggplot(data,aes(x = Price)) +  
  geom_histogram(aes(y =..density..),  
    bins= 25,  
    fill = "grey",  
    color ="black") +  
  geom_vline(xintercept = quantile(data$Price, 0.25), color = "dark red",lty = 2) +  
  geom_vline(xintercept = quantile(data$Price, 0.5), color = "dark red", ) +  
  geom_vline(xintercept = quantile(data$Price, 0.75), color = "dark red", lty = 2) +  
  labs(x = "Price", y ="Density") +  
  ggtitle("Price Distribution (quartiles)") +  
  geom_density()
```


Price Distribution (quartiles)




```
ggplot(data,aes(x = log(Price))) +
  geom_histogram(aes(y =..density..),
    bins= 25,
    fill = "grey",
    color ="black") +
  geom_vline(xintercept = quantile(data$LogPrice, 0.25), color = "dark red",lty = 2) +
  geom_vline(xintercept = quantile(data$LogPrice, 0.5), color = "dark red", ) +
  geom_vline(xintercept = quantile(data$LogPrice, 0.75), color = "dark red", lty = 2) +
  labs(x = "log(Price)", y ="Density") +
  ggtitle("log(Price) Distribution (quartiles)") +
  geom_density()
```




Descrittive variabile dipendente price


```
ggplot(data, aes(x = Price, fill = TypeName)) +  
  geom_density(size = 0.6, alpha = .3) +  
  labs(x = "Price", y = "Density", fill = "TypeName") +  
  ggtitle("Price Density Distribution For TypeName")
```



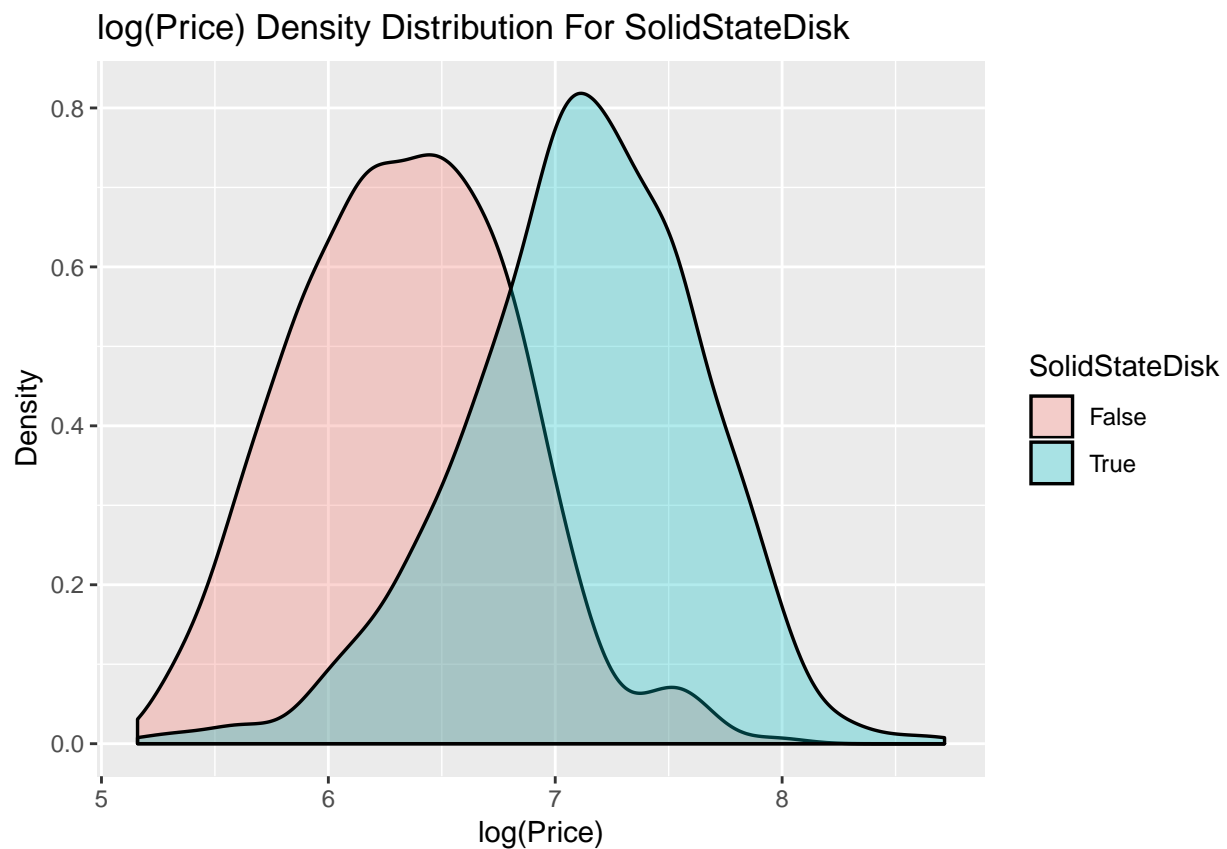
```
ggplot(data, aes(x = log(Price), fill = TypeName)) +  
  geom_density(size = 0.6, alpha = .3) +  
  labs(x = "log(Price)", y = "Density", fill = "TypeName") +  
  ggtitle("log(Price) Density Distribution For TypeName")
```



```
ggplot(data, aes(x = Price, fill = SolidStateDisk)) +  
  geom_density(size = 0.6, alpha = .3) +  
  labs(x = "Price", y = "Density", fill = "SolidStateDisk") +  
  ggtitle("Price Density Distribution For SolidStateDisk")
```



```
ggplot(data, aes(x = log(Price), fill = SolidStateDisk)) +  
  geom_density(size = 0.6, alpha = .3) +  
  labs(x = "log(Price)", y = "Density", fill = "SolidStateDisk") +  
  ggtitle("log(Price) Density Distribution For SolidStateDisk")
```



```
library(psych)
```

```
##
## Attaching package: 'psych'
## The following objects are masked from 'package:ggplot2':
##
##   %+%, alpha
```


```
describe(data$Price)
```

```
##   vars    n   mean    sd median trimmed   mad min  max range skew
## X1      1 1303 1123.69 699.01   977 1038.47 619.73 174 6099  5925 1.52
##   kurtosis    se
## X1        4.34 19.36
```

```
library(nortest)
```


```
# NORMALITA'
```

```
boxplot(data$Price)
```



```
qqnorm(data$Price);qqline(data$Price)
```

Normal Q-Q Plot



```
shapiro.test(data$Price)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  data$Price
## W = 0.89382, p-value < 2.2e-16
```

```
ad.test(data$Price)
```


```
##
##  Anderson-Darling normality test
```

```
##
## data: data$Price
## A = 28.319, p-value < 2.2e-16
#wilcox.test(data$Price, conf.int = TRUE, mu = ) #worth it?
#if(!require(Envstats)) install.packages("EnvStats")
library(EnvStats)

##
## Attaching package: 'EnvStats'
## The following objects are masked from 'package:stats':
##
## predict, predict.lm
## The following object is masked from 'package:base':
##
## print.default
varTest(sample(data$Price), sigma.squared = (sd(data$Price)*sd(data$Price)))


##
## Chi-Squared Test on Variance
##
## data: sample(data$Price)
## Chi-Squared = 1302, df = 1302, p-value = 0.9896
## alternative hypothesis: true variance is not equal to 488613.6
## 95 percent confidence interval:
## 453149.5 528432.0
## sample estimates:
## variance
## 488613.6

Trying with the log correction:
# Correzione NORMALITA'
library(nortest)
boxplot(data$LogPrice)
```



```
qqnorm(data$LogPrice);qqline(data$LogPrice)
```


Normal Q-Q Plot



```
shapiro.test(data$LogPrice) #better than before, but still not normal according to shapiro
```

```
##  
##  Shapiro-Wilk normality test  
##  
## data:  data$LogPrice  
## W = 0.99252, p-value = 3.628e-06
```

```
ad.test(data$LogPrice)
```

```
##  
##  Anderson-Darling normality test  
##  
## data:  data$LogPrice  
## A = 2.5942, p-value = 1.515e-06
```

T-test

```
# One sample  
ref <- mean(data$Price)  
Apple<-data$Price[data$Company=="Apple"]  
t.test(Apple,mu=ref,alternative = "greater")
```

```
##  
##  One Sample t-test  
##  
## data:  Apple  
## t = 3.5944, df = 20, p-value = 0.000906  
## alternative hypothesis: true mean is greater than 1123.687  
## 95 percent confidence interval:  
##  1352.823      Inf  
## sample estimates:
```

```
## mean of x
## 1564.199

# Wilcoxon Signed Rank Test
wilcox.test(Apple, mu=ref, conf.int = TRUE)

##
## Wilcoxon signed rank test
##
## data: Apple
## V = 206, p-value = 0.0008516
## alternative hypothesis: true location is not equal to 1123.687
## 95 percent confidence interval:
## 1234.50 1829.26
## sample estimates:
## (pseudo)median
## 1514.275
```

```
#Two sample
Other <-data$Price[data$Company!="Apple"]
wilcox.test(Apple, Other, alternative = "g")
```

```
##
## Wilcoxon rank sum test with continuity correction
##
## data: Apple and Other
## W = 19689, p-value = 0.0001358
## alternative hypothesis: true location shift is greater than 0
```

```
# F test sulla varianza
var.test(Apple, Other, alternative = "two.sided")
```

```
##
## F test to compare two variances
##
## data: Apple and Other
## F = 0.64574, num df = 20, denom df = 1281, p-value = 0.2401
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
## 0.3755878 1.3509884
## sample estimates:
## ratio of variances
## 0.6457382
```

Variabili qualitative: tabella di contingenza e chi quadro

```
b<-data
b.table<-table(b$SolidStateDisk,b$TypeName)
b.table
```

```
##
##      2 in 1 Convertible Gaming Netbook Notebook Ultrabook Workstation
## False      26      27      13      372      16      6
## True       95     178      12     355     180     23
```

```
prop.table(b.table,2)
```

```
##
##      2 in 1 Convertible      Gaming      Netbook      Notebook      Ultrabook
```

```
## False      0.21487603 0.13170732 0.52000000 0.51169188 0.08163265
## True       0.78512397 0.86829268 0.48000000 0.48830812 0.91836735
##
## Workstation
## False 0.20689655
## True  0.79310345
```

```
# chi square test
```

```
chisq.test(b.table)
```

```
##
## Pearson's Chi-squared test
##
## data: b.table
## X-squared = 203.18, df = 5, p-value < 2.2e-16
```

```
chi=chisq.test(b.table)
chi_norm=chi$statistic/(nrow(b)*min(nrow(b.table)-1,ncol(b.table)-1))
chi_norm
```

```
## X-squared
## 0.1559288
```

```
summary(b.table)
```

```
## Number of cases in table: 1303
## Number of factors: 2
## Test for independence of all factors:
## Chisq = 203.18, df = 5, p-value = 5.944e-42
```

Correlazione per variabili quantitative

```
# seleziona solo variabili quantitative
nums <- sapply(data, is.numeric)
var_numeric <- data[,nums]
head(var_numeric)
```

```
## X Inches Ram Weight Price Frequenza Pixel LogPrice
## 1 1 13.3 8 1.37 1339.69 2.3 4096000 7.200194
## 2 2 13.3 8 1.34 898.94 1.8 1296000 6.801216
## 3 3 15.6 8 1.86 575.00 2.5 2073600 6.354370
## 4 4 15.4 16 1.83 2537.45 2.7 5184000 7.838915
## 5 5 13.3 8 1.37 1803.60 3.1 4096000 7.497540
## 6 6 15.6 4 2.10 400.00 3.0 1049088 5.991465
```

```
var_numeric$X=NULL
```

```
# Matrice di correlazione
R<-cor(var_numeric)
R
```

```
## Inches Ram Weight Price Frequenza
## Inches 1.0000000 0.2379928 0.82763110 0.06819667 0.3078698
## Ram 0.23799280 1.0000000 0.38387409 0.74300714 0.3680005
## Weight 0.82763110 0.3838741 1.00000000 0.21036980 0.3204336
## Price 0.06819667 0.7430071 0.21036980 1.00000000 0.4302931
## Frequenza 0.30786980 0.3680005 0.32043359 0.43029310 1.0000000
## Pixel -0.08639917 0.3963585 -0.04403379 0.51548639 0.1352935
```

```
## LogPrice    0.04432871 0.6848033 0.15167383 0.92758068 0.5041461
##              Pixel    LogPrice
## Inches      -0.08639917 0.04432871
## Ram          0.39635848 0.68480333
## Weight       -0.04403379 0.15167383
## Price        0.51548639 0.92758068
## Frequenza    0.13529350 0.50414608
## Pixel        1.00000000 0.48490475
## LogPrice     0.48490475 1.00000000
```

```
# Test di correlazione. (Spearsman's o Kendall tau)
cor.test(var_numeric$Inches, var_numeric$Weight)
```

```
##
## Pearson's product-moment correlation
##
## data: var_numeric$Inches and var_numeric$Weight
## t = 53.187, df = 1301, p-value < 2.2e-16
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
##  0.8097181 0.8440031
## sample estimates:
##          cor
## 0.8276311
```

```
#corrgram(var_numeric)
```


```
# Correlazione come grafo
library(qgraph)
```

```
## Registered S3 methods overwritten by 'huge':
##   method      from
##   plot.sim    BDgraph
##   print.sim   BDgraph
```

```
detcor=cor(as.matrix(var_numeric), method="pearson")
round(detcor, 2)
```

```
##           Inches  Ram Weight Price Frequenza Pixel LogPrice
## Inches      1.00 0.24  0.83 0.07      0.31 -0.09  0.04
## Ram          0.24 1.00  0.38 0.74      0.37  0.40  0.68
## Weight       0.83 0.38  1.00 0.21      0.32 -0.04  0.15
## Price        0.07 0.74  0.21 1.00      0.43  0.52  0.93
## Frequenza    0.31 0.37  0.32 0.43      1.00  0.14  0.50
## Pixel       -0.09 0.40 -0.04 0.52      0.14  1.00  0.48
## LogPrice     0.04 0.68  0.15 0.93      0.50  0.48  1.00
```


```
# plot corr matrix: green positive red negative
qgraph(detcor, shape="circle", posCol="darkgreen", negCol="darkred")
```



Boxplot di confronto (pre-anova)


```
boxplot(data$Price~data$Company,  
        main="Boxplot Prezzo per compagnia",  
        col= rainbow(6),  
        horizontal = F)
```

Boxplot Prezzo per compagnia



```
boxplot(data$Price~data$SolidStateDisk,
        main="Prezzo vs ssd",
        col= rainbow(2),
        horizontal = F)
```

Prezzo vs ssd



ANOVA

A una via

```
lmA = lm(Price ~ SolidStateDisk, data=data)
summary(lmA)

##
## Call:
## lm(formula = Price ~ SolidStateDisk, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1214.8  -343.3   -96.8    261.1   4710.2
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      637.86      27.98   22.80  <2e-16 ***
## SolidStateDiskTrue  750.93      34.78   21.59  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 600 on 1301 degrees of freedom
## Multiple R-squared:  0.2638, Adjusted R-squared:  0.2632
## F-statistic: 466.2 on 1 and 1301 DF,  p-value: < 2.2e-16

drop1(lmA, test = 'F')

## Single term deletions
##
## Model:
## Price ~ SolidStateDisk
##              Df Sum of Sq      RSS   AIC F value    Pr(>F)
## <none>                 468355897 16672
## SolidStateDisk  1 167819064 636174961 17069  466.17 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

anova(lmA)

## Analysis of Variance Table
##
## Response: Price
##              Df      Sum Sq  Mean Sq F value    Pr(>F)
## SolidStateDisk    1 167819064 167819064  466.17 < 2.2e-16 ***
## Residuals      1301 468355897    359997
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

library(lsmeans)

## Loading required package: emmeans

## The 'lsmeans' package is now basically a front end for 'emmeans'.
## Users are encouraged to switch the rest of the way.
## See help('transition') for more information, including how to
## convert old 'lsmeans' objects and scripts to work with 'emmeans'.


ls_SolidStateDisk = lsmeans(lmA, pairwise ~ SolidStateDisk, adjust = 'tukey')
ls_SolidStateDisk$contrasts
```

```
## contrast      estimate    SE    df t.ratio p.value
## False - True      -751 34.8 1301 -21.591 <.0001
```

```
ls_SolidStateDisk$lsmeans
```

```
## SolidStateDisk lsmean    SE    df lower.CL upper.CL
## False           638 28.0 1301     583     693
## True            1389 20.7 1301    1348    1429
##
## Confidence level used: 0.95
```

```
plot(ls_SolidStateDisk$lsmeans, alpha = .05)
```



```
lmB = lm(Price ~ Company, data=data)
summary(lmB)
```

```
##
## Call:
## lm(formula = Price ~ Company, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2317.1  -452.8  -127.4   288.5  3812.6
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    626.78     63.43   9.881  < 2e-16 ***
## CompanyApple    937.42    154.14   6.082 1.57e-09 ***
## CompanyAsus    477.39     81.53   5.856 6.03e-09 ***
```



```
## CompanyChuwi      -312.48      377.06    -0.829 0.407416
## CompanyDell        559.29       73.62     7.597 5.80e-14 ***
## CompanyFujitsu    102.22      377.06     0.271 0.786352
## CompanyGoogle     1050.89     377.06     2.787 0.005397 **
## CompanyHP          441.00       74.41     5.927 3.96e-09 ***
## CompanyHuawei       797.22     459.62     1.735 0.083065 .
## CompanyLenovo      459.61       73.62     6.243 5.81e-10 ***
## CompanyLG          1472.22     377.06     3.904 9.93e-05 ***
## CompanyMediacom   -331.78      251.46    -1.319 0.187270
## CompanyMicrosoft  985.53      270.37     3.645 0.000278 ***
## CompanyMSI         1102.13     108.16    10.190 < 2e-16 ***
## CompanyRazer       2719.37     251.46    10.814 < 2e-16 ***
## CompanySamsung     786.67      223.77     3.515 0.000454 ***
## CompanyToshiba     641.04      112.51     5.698 1.50e-08 ***
## CompanyVero        -409.35      328.08    -1.248 0.212365
## CompanyXiaomi       506.69      328.08     1.544 0.122740
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 643.8 on 1284 degrees of freedom
## Multiple R-squared:  0.1635, Adjusted R-squared:  0.1518
## F-statistic: 13.94 on 18 and 1284 DF,  p-value: < 2.2e-16
```

```
drop1(lmB, test = 'F')
```

```
## Single term deletions
##
## Model:
## Price ~ Company
##           Df Sum of Sq      RSS   AIC F value    Pr(>F)
## <none>                 532160971 16873
## Company 18 104013991 636174961 17069   13.943 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(lmB)
```


```
## Analysis of Variance Table
##
## Response: Price
##           Df      Sum Sq Mean Sq F value    Pr(>F)
## Company    18 104013991 5778555   13.943 < 2.2e-16 ***
## Residuals 1284 532160971  414456
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
ls_Company = lsmeans(lmB, pairwise ~ Company, adjust = 'tukey')
#ls_Company$contrasts #too long to be printed
ls_Company$lsmeans
```

```
## Company  lsmean    SE    df lower.CL upper.CL
## Acer      627    63.4  1284   502.331     751
## Apple    1564   140.5  1284  1288.594    1840
## Asus     1104    51.2  1284  1003.692    1205
## Chuwi     314   371.7  1284  -414.885    1043
## Dell     1186    37.4  1284  1112.783    1259
```

```
## Fujitsu      729 371.7 1284   -0.182    1458
## Google      1678 371.7 1284  948.485    2407
## HP          1068  38.9 1284  991.475    1144
## Huawei      1424 455.2 1284  530.938    2317
## Lenovo      1086  37.4 1284 1013.099    1160
## LG          2099 371.7 1284 1369.818    2828
## Mediacom     295 243.3 1284 -182.362     772
## Microsoft   1612 262.8 1284 1096.699    2128
## MSI          1729  87.6 1284 1557.038    1901
## Razer        3346 243.3 1284 2868.781    3824
## Samsung      1413 214.6 1284  992.451    1834
## Toshiba      1268  92.9 1284 1085.517    1450
## Vero         217 321.9 1284 -414.065     849
## Xiaomi      1133 321.9 1284  501.972    1765
##
## Confidence level used: 0.95
```

```
plot(ls_Company$lsmeans, alpha = .05)
```



```
lmC = lm(Price ~ TypeName, data=data)
summary(lmC)
```

```
##
## Call:
## lm(formula = Price ~ TypeName, data = data)
##
## Residuals:
```

##	Min	1Q	Median	3Q	Max

```
## -1049.2 -381.7 -98.1 267.6 4367.6
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1282.40     50.01  25.642 < 2e-16 ***
## TypeNameGaming    448.98     63.07   7.119 1.79e-12 ***
## TypeNameNetbook  -646.17    120.86  -5.347 1.06e-07 ***
## TypeNameNotebook -500.32     54.01  -9.263 < 2e-16 ***
## TypeNameUltrabook  265.83     63.60   4.180 3.12e-05 ***
## TypeNameWorkstation 997.96    113.74   8.774 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 550.1 on 1297 degrees of freedom
## Multiple R-squared:  0.383, Adjusted R-squared:  0.3806
## F-statistic: 161 on 5 and 1297 DF, p-value: < 2.2e-16
```

```
drop1(lmC, test = 'F')
```

```
## Single term deletions
##
## Model:
## Price ~ TypeName
##              Df Sum of Sq      RSS   AIC F value    Pr(>F)
## <none>                 392518380 16450
## TypeName  5 243656581 636174961 17069  161.02 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(lmC)
```

```
## Analysis of Variance Table
##
## Response: Price
##              Df      Sum Sq Mean Sq F value    Pr(>F)
## TypeName      5 243656581 48731316  161.02 < 2.2e-16 ***
## Residuals 1297 392518380   302636
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
ls_TypeName = lsmeans(lmC, pairwise ~ TypeName, adjust = 'tukey')
ls_TypeName$contrasts
```


```
## contrast              estimate      SE    df t.ratio p.value
## 2 in 1 Convertible - Gaming      -449  63.1 1297  -7.119 <.0001
## 2 in 1 Convertible - Netbook      646 120.9 1297   5.347 <.0001
## 2 in 1 Convertible - Notebook     500  54.0 1297   9.263 <.0001
## 2 in 1 Convertible - Ultrabook    -266  63.6 1297  -4.180 0.0004
## 2 in 1 Convertible - Workstation  -998 113.7 1297  -8.774 <.0001
## Gaming - Netbook                1095 116.5 1297   9.397 <.0001
## Gaming - Notebook                949  43.5 1297  21.821 <.0001
## Gaming - Ultrabook              183  55.0 1297   3.333 0.0114
## Gaming - Workstation            -549 109.1 1297  -5.030 <.0001
## Netbook - Notebook             -146 111.9 1297  -1.303 0.7833
## Netbook - Ultrabook            -912 116.8 1297  -7.806 <.0001
## Netbook - Workstation          -1644 150.1 1297 -10.951 <.0001
```

```
## Notebook - Ultrabook          -766  44.3 1297 -17.304 <.0001
## Notebook - Workstation        -1498 104.2 1297 -14.383 <.0001
## Ultrabook - Workstation       -732 109.5 1297  -6.689 <.0001
##
## P value adjustment: tukey method for comparing a family of 6 estimates
```


```
ls_TypeName$lsmeans
```

```
## TypeName      lsmean    SE   df lower.CL upper.CL
## 2 in 1 Convertible 1282  50.0 1297    1184    1381
## Gaming            1731  38.4 1297    1656    1807
## Netbook           636 110.0 1297     420     852
## Notebook          782  20.4 1297     742     822
## Ultrabook         1548  39.3 1297    1471    1625
## Workstation       2280 102.2 1297    2080    2481
##
## Confidence level used: 0.95
```





```
plot(ls_TypeName$lsmeans, alpha = .05)
```



```
library(coefplot)
#library(forestmodel)
coefplot(lmC, intercept = FALSE)
```




```
par(mfrow = c(2,2))
plot(lmC)
```



```
##(not) normal distribution of residuals
par(mfrow=c(1,2))
boxplot(lmC$residuals)
qqnorm(lmC$residuals);qqline(lmC$residuals)
```

Normal Q-Q Plot



```
ad.test(lmC$residuals)
```

```
##
## Anderson-Darling normality test
##
## data:  lmC$residuals
## A = 22.667, p-value < 2.2e-16
```

```
shapiro.test(lmC$residuals)
```

```
##
## Shapiro-Wilk normality test
##
## data:  lmC$residuals
## W = 0.89641, p-value < 2.2e-16
```

```
##let's try again with the log correction
lmC_log = lm(log(Price) ~ TypeName, data=data)
summary(lmC_log)##R^2 increases
```

```
##
## Call:
## lm(formula = log(Price) ~ TypeName, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -1.40971 -0.33589 0.00698 0.33215 1.96853
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      7.02648    0.04379 160.456 < 2e-16 ***
## TypeNameGaming      0.33865    0.05522   6.133 1.15e-09 ***
## TypeNameNetbook    -0.91149    0.10583  -8.613 < 2e-16 ***
## TypeNameNotebook   -0.49823    0.04729 -10.534 < 2e-16 ***
## TypeNameUltrabook   0.26648    0.05569   4.785 1.91e-06 ***
## TypeNameWorkstation 0.66479    0.09959   6.675 3.65e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4817 on 1297 degrees of freedom
## Multiple R-squared:  0.4061, Adjusted R-squared:  0.4038
## F-statistic: 177.4 on 5 and 1297 DF,  p-value: < 2.2e-16
```

```
drop1(lmC_log, test = 'F')
```

```
## Single term deletions
##
## Model:
## log(Price) ~ TypeName
##              Df Sum of Sq    RSS      AIC F value    Pr(>F)
## <none>                    300.95 -1897.5
## TypeName  5      205.76 506.71 -1228.7  177.36 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(lmC_log)
```

```
## Analysis of Variance Table
##
## Response: log(Price)
##              Df Sum Sq Mean Sq F value    Pr(>F)
## TypeName      5  205.76   41.152  177.36 < 2.2e-16 ***
## Residuals 1297  300.95    0.232
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
ls_TypeName_log = lsmeans(lmC_log, pairwise ~ TypeName, adjust = 'tukey')
ls_TypeName_log$contrasts
```


```
## contrast              estimate      SE    df t.ratio p.value
## 2 in 1 Convertible - Gaming      -0.3387 0.0552 1297  -6.133 <.0001
## 2 in 1 Convertible - Netbook       0.9115 0.1058 1297   8.613 <.0001
## 2 in 1 Convertible - Notebook      0.4982 0.0473 1297  10.534 <.0001
## 2 in 1 Convertible - Ultrabook    -0.2665 0.0557 1297  -4.785 <.0001
## 2 in 1 Convertible - Workstation  -0.6648 0.0996 1297  -6.675 <.0001
## Gaming - Netbook                  1.2501 0.1020 1297  12.251 <.0001
## Gaming - Notebook                  0.8369 0.0381 1297  21.970 <.0001
## Gaming - Ultrabook                 0.0722 0.0481 1297   1.500 0.6644
## Gaming - Workstation              -0.3261 0.0956 1297  -3.413 0.0087
## Netbook - Notebook                -0.4133 0.0980 1297  -4.218 0.0004
## Netbook - Ultrabook               -1.1780 0.1023 1297 -11.515 <.0001
## Netbook - Workstation             -1.5763 0.1315 1297 -11.990 <.0001
```

```
## Notebook - Ultrabook          -0.7647 0.0388 1297 -19.725 <.0001
## Notebook - Workstation        -1.1630 0.0912 1297 -12.750 <.0001
## Ultrabook - Workstation       -0.3983 0.0958 1297  -4.156 0.0005
##
## Results are given on the log (not the response) scale.
## P value adjustment: tukey method for comparing a family of 6 estimates
```


```
ls_TypeName_log$lsmeans
```

```
## TypeName      lsmean      SE   df lower.CL upper.CL
## 2 in 1 Convertible  7.03 0.0438 1297    6.94    7.11
## Gaming            7.37 0.0336 1297    7.30    7.43
## Netbook           6.11 0.0963 1297    5.93    6.30
## Notebook          6.53 0.0179 1297    6.49    6.56
## Ultrabook         7.29 0.0344 1297    7.23    7.36
## Workstation       7.69 0.0894 1297    7.52    7.87
##
## Results are given on the log (not the response) scale.
## Confidence level used: 0.95
```

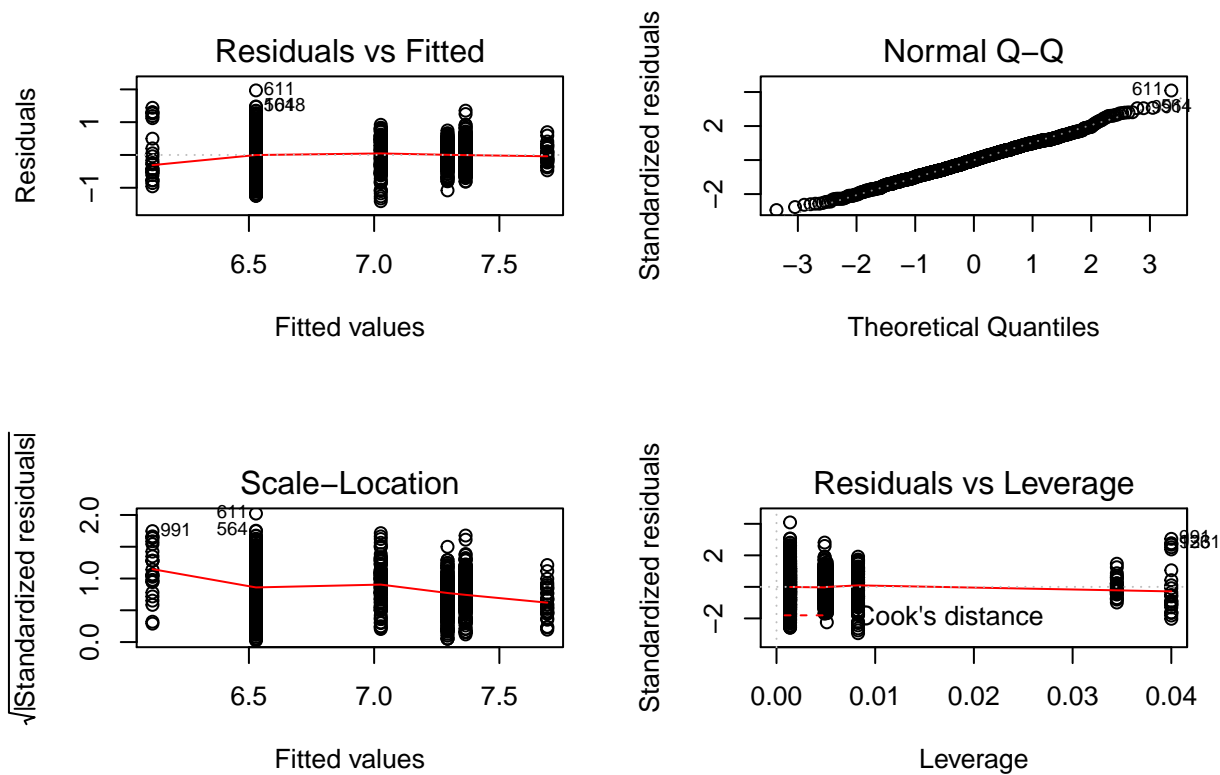
```
plot(ls_TypeName_log$lsmeans, alpha = .05)
```



```
coefplot(lmC_log, intercept = FALSE)
```

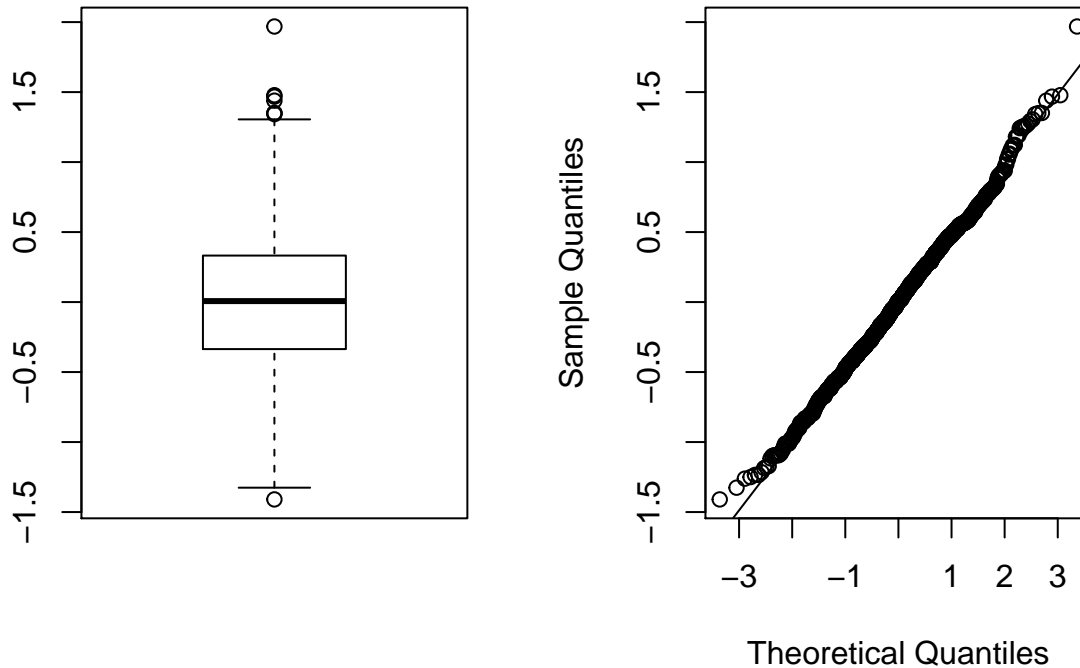



```
par(mfrow = c(2,2))
plot(lmC_log)
```



```
#(not) normal distribution of residuals
par(mfrow=c(1,2))
boxplot(lmC_log$residuals)
qqnorm(lmC_log$residuals);qqline(lmC_log$residuals)
```

Normal Q-Q Plot



```
ad.test(lmC_log$residuals) #normal now!
```

```
##
## Anderson-Darling normality test
##
## data: lmC_log$residuals
## A = 0.51757, p-value = 0.1886
```

```
shapiro.test(lmC_log$residuals) #borderline now!
```

```
##
## Shapiro-Wilk normality test
##
## data: lmC_log$residuals
## W = 0.99764, p-value = 0.05462
```

A due vie

```
# Con interazione
lmC = lm(Price ~ Company*TypeName , data=data)
drop1(lmC, test="F")
```

```
## Single term deletions
##
## Model:
## Price ~ Company * TypeName
```

```
##              Df Sum of Sq      RSS      AIC F value    Pr(>F)
## <none>                        320739568 16273
## Company:TypeName 25  29159364 349898932 16336  4.5602 1.181e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#summary(lmC) #FIXME: too long to be printed
```

```
lmC = lm(Price ~ Company+TypeName , data=data)
# type I effects A, B/A C/A,B
anova(lmC)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: Price
```

```
##              Df      Sum Sq Mean Sq F value    Pr(>F)
## Company      18 104013991  5778555  21.123 < 2.2e-16 ***
## TypeName      5 182262038 36452408 133.246 < 2.2e-16 ***
## Residuals 1279 349898932   273572
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# type III effects A/B,C , B/A,C C/A,B
```

```
drop1(lmC, test="F")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## Price ~ Company + TypeName
```

```
##              Df Sum of Sq      RSS      AIC F value    Pr(>F)
## <none>                        349898932 16336
## Company  18  42619448 392518380 16450   8.6549 < 2.2e-16 ***
## TypeName  5 182262038 532160971 16873 133.2460 < 2.2e-16 ***
## ---
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(lmC)
```

```
##
```

```
## Call:
```

```
## lm(formula = Price ~ Company + TypeName, data = data)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
## -2147.6  -343.2   -81.9    243.1   4081.9
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      991.52      69.88  14.189 < 2e-16 ***
## CompanyApple      383.70     132.62   2.893  0.00388 **
## CompanyAsus       168.81      67.79   2.490  0.01290 *
## CompanyChuwi     -180.68     306.47  -0.590  0.55559
## CompanyDell       350.73      60.52   5.796 8.56e-09 ***
## CompanyFujitsu    234.02     306.47   0.764  0.44525
## CompanyGoogle     497.17     309.44   1.607  0.10837
## CompanyHP         337.48      60.85   5.546 3.55e-08 ***
```

```


## CompanyHuawei      243.50      375.96      0.648      0.51731
## CompanyLenovo      322.12       60.24      5.348 1.05e-07 ***
## CompanyLG          918.50     309.44      2.968      0.00305 **
## CompanyMediacom    -270.91     204.43     -1.325      0.18534
## CompanyMicrosoft   431.81     223.95      1.928      0.05406 .
## CompanyMSI         311.10      98.62      3.155      0.00165 **
## CompanyRazer       1996.14     207.24      9.632 < 2e-16 ***
## CompanySamsung     438.88     183.82      2.388      0.01710 *
## CompanyToshiba     601.45      92.12      6.529 9.52e-11 ***
## CompanyVero        -277.55     266.70     -1.041      0.29821
## CompanyXiaomi       295.72     267.40      1.106      0.26896
## TypeNameGaming      426.29      65.51      6.507 1.10e-10 ***
## TypeNameNetbook    -600.94     115.75     -5.192 2.42e-07 ***
## TypeNameNotebook   -496.54      51.98     -9.552 < 2e-16 ***
## TypeNameUltrabook   188.98      63.81      2.962      0.00312 **
## TypeNameWorkstation 948.46     109.22      8.684 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 523 on 1279 degrees of freedom
## Multiple R-squared:  0.45, Adjusted R-squared:  0.4401
## F-statistic: 45.5 on 23 and 1279 DF, p-value: < 2.2e-16

# contrasti
library(lsmeans)
ls=lsmeans(lmC, #FIXME: @Andrea, c'era lmB ma credo tu volessi scrivere lmC, in case check it
            pairwise ~ TypeName ,
            adjust="tukey")
ls$lsmeans

##   TypeName      lsmean      SE    df lower.CL upper.CL
## 2 in 1 Convertible  1350  68.9 1279     1214     1485
## Gaming             1776  64.4 1279     1649     1902
## Netbook             749 115.9 1279      521      976
## Notebook            853  52.0 1279      751      955
## Ultrabook          1538  55.5 1279     1430     1647
## Workstation        2298 110.1 1279     2082     2514
##
## Results are averaged over the levels of: Company
## Confidence level used: 0.95

# plot lsmeans and 95% confid int
plot(ls$lsmeans, alpha = .05)

```



```
# contrasts between predicted lsmeans
ls$contrasts
```

```
## contrast          estimate    SE   df t.ratio p.value
## 2 in 1 Convertible - Gaming      -426  65.5 1279  -6.507 <.0001
## 2 in 1 Convertible - Netbook       601 115.7 1279   5.192 <.0001
## 2 in 1 Convertible - Notebook      497  52.0 1279   9.552 <.0001
## 2 in 1 Convertible - Ultrabook    -189  63.8 1279  -2.962 0.0367
## 2 in 1 Convertible - Workstation  -948 109.2 1279  -8.684 <.0001
## Gaming - Netbook                1027 114.5 1279   8.972 <.0001
## Gaming - Notebook                923  49.4 1279  18.671 <.0001
## Gaming - Ultrabook              237  61.1 1279   3.882 0.0015
## Gaming - Workstation            -522 108.3 1279  -4.820 <.0001
## Netbook - Notebook             -104 107.0 1279  -0.975 0.9258
## Netbook - Ultrabook            -790 113.3 1279  -6.969 <.0001
## Netbook - Workstation         -1549 143.8 1279 -10.774 <.0001
## Notebook - Ultrabook           -686  46.5 1279 -14.754 <.0001
## Notebook - Workstation        -1445  99.8 1279 -14.475 <.0001
## Ultrabook - Workstation        -759 106.4 1279  -7.138 <.0001
##
```


```
## Results are averaged over the levels of: Company
## P value adjustment: tukey method for comparing a family of 6 estimates
```

```
# if at least one contrast is significant, the variable
# is significant in the anova table # drop1 effects
```

```
# contrast among predicted lsmeans and overall lsmean
c= contrast(ls, method = "eff")
```

c

```
## $lsmeans
## contrast estimate SE df t.ratio p.value
## 2 in 1 Convertible effect -77.7 47.9 1279 -1.623 0.1048
## Gaming effect 348.6 46.0 1279 7.583 <.0001
## Netbook effect -678.6 90.2 1279 -7.521 <.0001
## Notebook effect -574.2 31.8 1279 -18.032 <.0001
## Ultrabook effect 111.3 43.8 1279 2.542 0.0134
## Workstation effect 870.7 84.6 1279 10.287 <.0001
##
## Results are averaged over the levels of: Company
## P value adjustment: fdr method for 6 tests
##
## $contrasts
## contrast estimate SE df t.ratio
## 2 in 1 Convertible - Gaming effect -150.6 71.6 1279 -2.103
## 2 in 1 Convertible - Netbook effect 876.6 121.9 1279 7.192
## 2 in 1 Convertible - Notebook effect 772.2 51.2 1279 15.077
## 2 in 1 Convertible - Ultrabook effect 86.7 57.9 1279 1.498
## 2 in 1 Convertible - Workstation effect -672.8 74.0 1279 -9.093
## Gaming - Netbook effect 1302.9 123.6 1279 10.544
## Gaming - Notebook effect 1198.5 55.4 1279 21.649
## Gaming - Ultrabook effect 513.0 60.9 1279 8.416
## Gaming - Workstation effect -246.5 77.4 1279 -3.186
## Netbook - Notebook effect 171.2 107.0 1279 1.600
## Netbook - Ultrabook effect -514.3 110.5 1279 -4.655
## Netbook - Workstation effect -1273.7 119.6 1279 -10.649
## Notebook - Ultrabook effect -409.9 55.3 1279 -7.416
## Notebook - Workstation effect -1169.3 71.6 1279 -16.325
## Ultrabook - Workstation effect -483.8 84.4 1279 -5.730
## p.value
## 0.0411
## <.0001
## <.0001
## 0.1345
## <.0001
## <.0001
## <.0001
## <.0001
## <.0001
## 0.0018
## 0.1177
## <.0001
## <.0001
## <.0001
## <.0001
## <.0001
##
## Results are averaged over the levels of: Company
## P value adjustment: fdr method for 15 tests
library(coefplot)
coefplot(lmC, intercept=FALSE) #FIXME: @Andrea, same goes here
```



ANOVA k way

```
lmK = lm(Price ~ Company+TypeName+SolidStateDisk , data=data)
summary(lmK)
```


```
##
## Call:
## lm(formula = Price ~ Company + TypeName + SolidStateDisk, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2113.2  -301.6   -49.8    210.4   3862.4
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      689.90      67.44  10.230 < 2e-16 ***
## CompanyApple      514.90     122.54   4.202 2.83e-05 ***
## CompanyAsus       130.26      62.53   2.083 0.03744 *
## CompanyChuwi     -16.66     282.67  -0.059 0.95300
## CompanyDell       280.98      55.97   5.020 5.88e-07 ***
## CompanyFujitsu     84.49     282.64   0.299 0.76505
## CompanyGoogle     404.40     285.26   1.418 0.15653
## CompanyHP         272.28      56.25   4.840 1.45e-06 ***
## CompanyHuawei      150.73     346.56   0.435 0.66368
## CompanyLenovo     235.35      55.81   4.217 2.65e-05 ***
## CompanyLG         825.73     285.26   2.895 0.00386 **
## CompanyMediacom  -356.00     188.50  -1.889 0.05918 .
## CompanyMicrosoft  339.04     206.50   1.642 0.10087
```

```

## CompanyMSI          178.78      91.32      1.958  0.05047 .
## CompanyRazer        1868.90     191.19      9.775 < 2e-16 ***
## CompanySamsung      387.48     169.45      2.287  0.02238 *
## CompanyToshiba      436.72      85.60      5.102 3.87e-07 ***
## CompanyVero         -113.54     246.04     -0.461  0.64456
## CompanyXiaomi        96.18      246.80      0.390  0.69680
## TypeNameGaming       398.61      60.40      6.599 6.05e-11 ***
## TypeNameNetbook     -473.92     107.01     -4.429 1.03e-05 ***
## TypeNameNotebook    -358.93      48.77     -7.360 3.28e-13 ***
## TypeNameUltrabook    113.04      59.02      1.915  0.05570 .
## TypeNameWorkstation  946.96     100.66      9.407 < 2e-16 ***
## SolidStateDiskTrue   470.33      31.17     15.089 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 482.1 on 1278 degrees of freedom
## Multiple R-squared:  0.5332, Adjusted R-squared:  0.5244
## F-statistic: 60.82 on 24 and 1278 DF,  p-value: < 2.2e-16
drop1(lmK, test="F") # type III SS

## Single term deletions
##
## Model:
## Price ~ Company + TypeName + SolidStateDisk
##              Df Sum of Sq      RSS   AIC  F value    Pr(>F)
## <none>                296988657 16125
## Company             18  33990309 330978966 16230    8.1259 < 2.2e-16 ***
## TypeName              5 109128253 406116910 16523   93.9200 < 2.2e-16 ***
## SolidStateDisk       1  52910275 349898932 16336  227.6832 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
coefplot(lmK, intercept=FALSE)

```





Regressione lineare

```
lmA<-lm(Price ~ Frequenza , data=data)
summary(lmA)
```

```
##
## Call:
## lm(formula = Price ~ Frequenza, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1467.6  -453.8  -119.6   327.6  4618.2
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -241.84     81.32  -2.974   0.003 **
## Frequenza     594.02     34.55  17.194 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 631.2 on 1301 degrees of freedom
## Multiple R-squared:  0.1852, Adjusted R-squared:  0.1845
## F-statistic: 295.6 on 1 and 1301 DF, p-value: < 2.2e-16

plot(data$Frequenza,data$Price)
abline(lmA,col="red")
```



```
lmA<-lm(Price ~ Frequenza+Pixel+Ram , data=data)
summary(lmA)
```

```
##
## Call:
## lm(formula = Price ~ Frequenza + Pixel + Ram, data = data)
##
## Residuals:
```


	Min	1Q	Median	3Q	Max
	-1785.72	-257.23	-66.06	191.11	2791.53

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-4.076e+02	5.547e+01	-7.349	3.52e-13 ***
Frequenza	2.549e+02	2.474e+01	10.306	< 2e-16 ***
Pixel	1.329e-04	9.117e-06	14.575	< 2e-16 ***
Ram	7.839e+01	2.658e+00	29.488	< 2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 420.2 on 1299 degrees of freedom
## Multiple R-squared:  0.6395, Adjusted R-squared:  0.6386
## F-statistic: 768 on 3 and 1299 DF, p-value: < 2.2e-16
```

```
coefplot(lmA, intercept=FALSE)
```



ANCOVA

```
lmK = lm(Price ~ Company+TypeName+SolidStateDisk+ Frequenza+Pixel+Ram , data=data)
summary(lmK)
```

```
##
## Call:
## lm(formula = Price ~ Company + TypeName + SolidStateDisk + Frequenza +
##     Pixel + Ram, data = data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1838.5  -211.8   -28.2   169.3  1894.4
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -1.491e+02  6.691e+01  -2.229  0.02602 *
## CompanyApple    2.826e+02  9.043e+01   3.125  0.00182 **
## CompanyAsus     5.438e+01  4.587e+01   1.185  0.23609
## CompanyChuwi   -7.683e+01  2.082e+02  -0.369  0.71213
## CompanyDell     1.124e+02  4.132e+01   2.720  0.00662 **
## CompanyFujitsu  5.168e+01  2.071e+02   0.250  0.80294
## CompanyGoogle   3.062e+02  2.105e+02   1.455  0.14602
## CompanyHP       2.045e+02  4.134e+01   4.947 8.54e-07 ***
## CompanyHuawei    5.510e+01  2.539e+02   0.217  0.82822
## CompanyLenovo   1.260e+02  4.108e+01   3.066  0.00221 **
## CompanyLG       6.759e+02  2.090e+02   3.235  0.00125 **
## CompanyMediacom -1.108e+02  1.392e+02  -0.796  0.42603
```

```
## CompanyMicrosoft      2.369e+02  1.515e+02  1.564  0.11807
## CompanyMSI            2.046e+02  6.686e+01  3.061  0.00225 **
## CompanyRazer          1.085e+03  1.428e+02  7.594  5.95e-14 ***
## CompanySamsung        9.436e+01  1.246e+02  0.757  0.44896
## CompanyToshiba        2.871e+02  6.306e+01  4.553  5.79e-06 ***
## CompanyVero           1.440e+01  1.811e+02  0.080  0.93663
## CompanyXiaomi         -1.743e+01  1.808e+02 -0.096  0.92322
## TypeNameGaming        -2.977e+01  4.812e+01 -0.619  0.53621
## TypeNameNetbook       -1.142e+02  7.947e+01 -1.437  0.15105
## TypeNameNotebook      -2.440e+02  3.642e+01 -6.700  3.11e-11 ***
## TypeNameUltrabook     9.405e+01  4.338e+01  2.168  0.03034 *
## TypeNameWorkstation   7.172e+02  7.500e+01  9.562  < 2e-16 ***
## SolidStateDiskTrue    1.997e+02  2.432e+01  8.212  5.28e-16 ***
## Frequenza             1.701e+02  2.320e+01  7.335  3.94e-13 ***
## Pixel                 8.315e-05  8.292e-06  10.028  < 2e-16 ***
## Ram                   6.541e+01  2.578e+00  25.368  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 352.9 on 1275 degrees of freedom
## Multiple R-squared:  0.7504, Adjusted R-squared:  0.7452
## F-statistic: 142 on 27 and 1275 DF, p-value: < 2.2e-16
```

```
drop1(lmK, .~., test="F")
```

```
## Single term deletions
##
## Model:
## Price ~ Company + TypeName + SolidStateDisk + Frequenza + Pixel +
##      Ram
##
##           Df Sum of Sq      RSS   AIC  F value    Pr(>F)
## <none>                 158760389 15315
## Company          18  13404444 172164833 15384    5.9806 4.092e-14 ***
## TypeName          5   35077529 193837917 15565   56.3413 < 2.2e-16 ***
## SolidStateDisk    1    8397143 167157532 15380   67.4372 5.281e-16 ***
## Frequenza         1    6698755 165459144 15367   53.7975 3.940e-13 ***
## Pixel             1   12521049 171281438 15412  100.5562 < 2.2e-16 ***
## Ram               1    80130237 238890626 15845  643.5236 < 2.2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
ls=lsmeans(lmK,
            pairwise ~ Company ,
            adjust="tukey")
c= contrast(ls, method = "eff")
#c #FIXME: too long to be printed
```

```
data$LogPrice=NULL
data$Product=NULL
data$X=NULL
str(data)
```

```
## 'data.frame':   1303 obs. of  15 variables:
## $ Company      : Factor w/ 19 levels "Acer","Apple",...: 2 2 8 2 2 1 2 2 3 1 ...
## $ TypeName     : Factor w/ 6 levels "2 in 1 Convertible",...: 5 5 4 5 5 4 5 5 5 5 ...
## $ Inches       : num  13.3 13.3 15.6 15.4 13.3 15.6 15.4 13.3 14 14 ...
```

```
## $ ScreenResolution: Factor w/ 40 levels "1366x768","1440x900",...: 24 2 9 26 24 1 26 2 9 16 ...
## $ Cpu              : Factor w/ 118 levels "AMD A10-Series 9600P 2.4GHz",...: 55 53 64 75 57 15 74 53 ...
## $ Ram              : num 8 8 8 16 8 4 16 8 16 8 ...
## $ Memory           : Factor w/ 39 levels "1.0TB HDD","1.0TB Hybrid",...: 5 3 17 30 17 27 16 16 30 17 ...
## $ Gpu              : Factor w/ 110 levels "AMD FirePro W4190M",...: 59 52 54 10 60 18 61 52 98 62 ...
## $ OpSys            : Factor w/ 9 levels "Android","Chrome OS",...: 5 5 6 5 5 7 4 5 7 7 ...
## $ Weight           : num 1.37 1.34 1.86 1.83 1.37 2.1 2.04 1.34 1.3 1.6 ...
## $ Price            : num 1340 899 575 2537 1804 ...
## $ Frequenza        : num 2.3 1.8 2.5 2.7 3.1 3 2.2 1.8 1.8 1.6 ...
## $ Risoluzione      : Factor w/ 15 levels "1366x768","1440x900",...: 11 2 4 13 11 1 13 2 4 4 ...
## $ Pixel            : int 4096000 1296000 2073600 5184000 4096000 1049088 5184000 1296000 2073600 2073600 ...
## $ SolidStateDisk   : Factor w/ 2 levels "False","True": 2 1 2 2 2 1 1 1 2 2 ...
```

```
lm_full = lm(Price ~ ., data = data)
#summary(lm_full) #FIXME: wayyy too long to be printed, R^2 = 0.9586
anova(lm_full, test="F")
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: Price
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Company	18	104013991	5778555	114.5882	< 2.2e-16 ***
TypeName	5	182262038	36452408	722.8478	< 2.2e-16 ***
Inches	1	6163570	6163570	122.2230	< 2.2e-16 ***
ScreenResolution	36	108074619	3002073	59.5308	< 2.2e-16 ***
Cpu	110	95329933	866636	17.1853	< 2.2e-16 ***
Ram	1	34947028	34947028	692.9963	< 2.2e-16 ***
Memory	35	17134540	489558	9.7079	< 2.2e-16 ***
Gpu	88	34242874	389124	7.7163	< 2.2e-16 ***
OpSys	6	3526085	587681	11.6537	1.198e-12 ***
Weight	1	973	973	0.0193	0.8895
Residuals	1001	50479311	50429		

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
drop1(lm_full, test="F")
```

```
## Single term deletions
```

```
##
```

```
## Model:
```

```
## Price ~ Company + TypeName + Inches + ScreenResolution + Cpu +
##       Ram + Memory + Gpu + OpSys + Weight + Frequenza + Risoluzione +
##       Pixel + SolidStateDisk
```

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
<none>			50479311	14370		
Company	14	6197922	56677232	14493	8.7789	< 2.2e-16 ***
TypeName	5	3685931	54165241	14452	14.6183	7.50e-14 ***
Inches	1	210134	50689445	14373	4.1669	0.04148 *
ScreenResolution	23	5322877	55802188	14454	4.5892	8.09e-12 ***
Cpu	88	16408116	66887427	14560	3.6974	< 2.2e-16 ***
Ram	1	4481351	54960662	14479	88.8648	< 2.2e-16 ***
Memory	34	10507055	60986365	14548	6.1281	< 2.2e-16 ***
Gpu	88	30459868	80939179	14809	6.8638	< 2.2e-16 ***
OpSys	6	3518495	53997806	14446	11.6286	1.28e-12 ***
Weight	1	973	50480284	14368	0.0193	0.88953

```
## Frequenza      0      0 50479311 14370
## Risoluzione    0      0 50479311 14370
## Pixel          0      0 50479311 14370
## SolidStateDisk 0      0 50479311 14370
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
#coefplot(lm_full, intercept=FALSE) #meglio di no ahah
```

```
par(mfrow=c(2,2))
plot(lm_full)
```

```
## Warning: not plotting observations with leverage one:
```

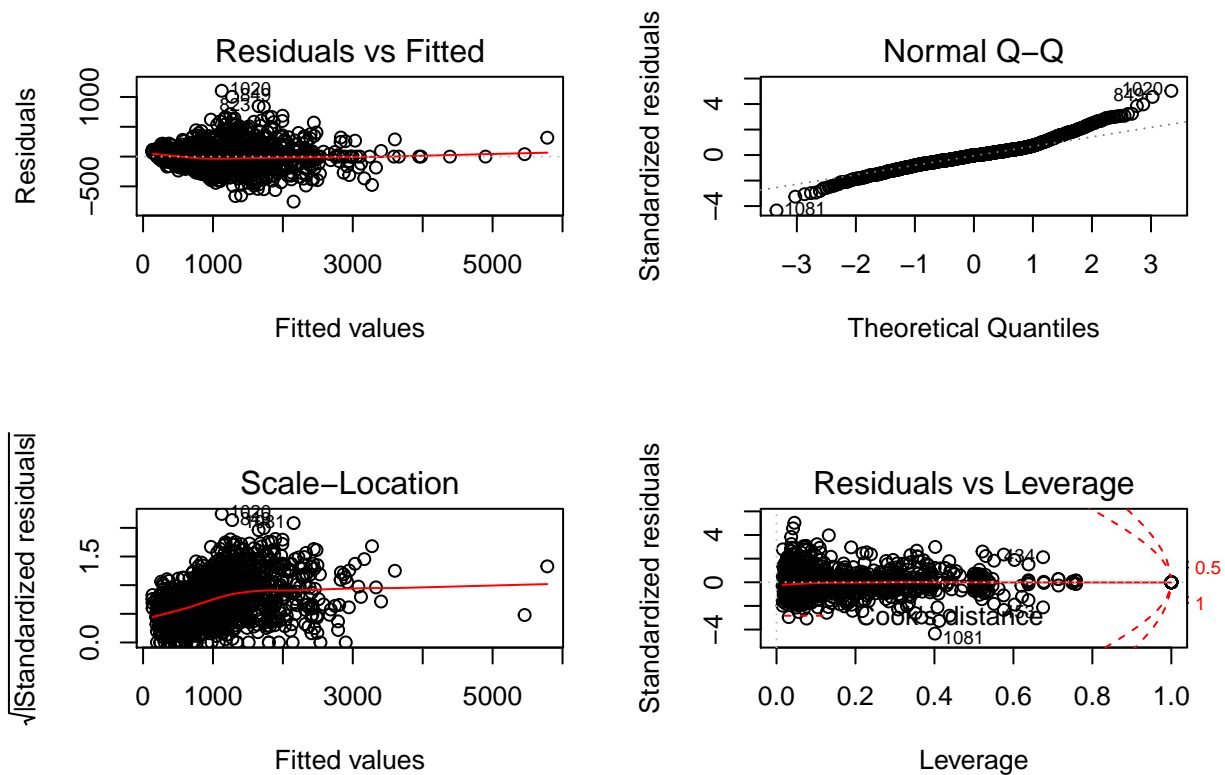
```
## 13, 15, 18, 29, 34, 46, 84, 128, 160, 173, 178, 179, 205, 232, 267, 271, 299, 303, 324, 348, 388, 4
```

```
## Warning: not plotting observations with leverage one:
```



```
## 13, 15, 18, 29, 34, 46, 84, 128, 160, 173, 178, 179, 205, 232, 267, 271, 299, 303, 324, 348, 388, 4
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```



```
par(mfrow=c(1,1))
par(mfrow=c(1,2))
boxplot(lm_full$residuals)
qqnorm(lm_full$residuals);qqline(lm_full$residuals) # probably the correction would work pretty fine here
```



```
#tests
ad.test(lm_full$residuals)

##
## Anderson-Darling normality test
##
## data:  lm_full$residuals
## A = 19.821, p-value < 2.2e-16


shapiro.test(lm_full$residuals)

##
## Shapiro-Wilk normality test
##
## data:  lm_full$residuals
## W = 0.94917, p-value < 2.2e-16

library(MASS)

##
## Attaching package: 'MASS'
## The following object is masked from 'package:EnvStats':
##
##      boxcox

boxcoxreg1<-boxcox(lm_full)
```



```
which.max(boxcoxreg1$y)
```

```
## [1] 48
```

```
lambda=boxcoxreg1$x[which.max(boxcoxreg1$y)]
lambda
```

```
## [1] -0.1010101
```

```
lm_full_t = lm(log(Price) ~ ., data = data)
```

```
par(mfrow=c(2,2))
```

```
plot(lm_full_t) #quite better
```

```
## Warning: not plotting observations with leverage one:
```

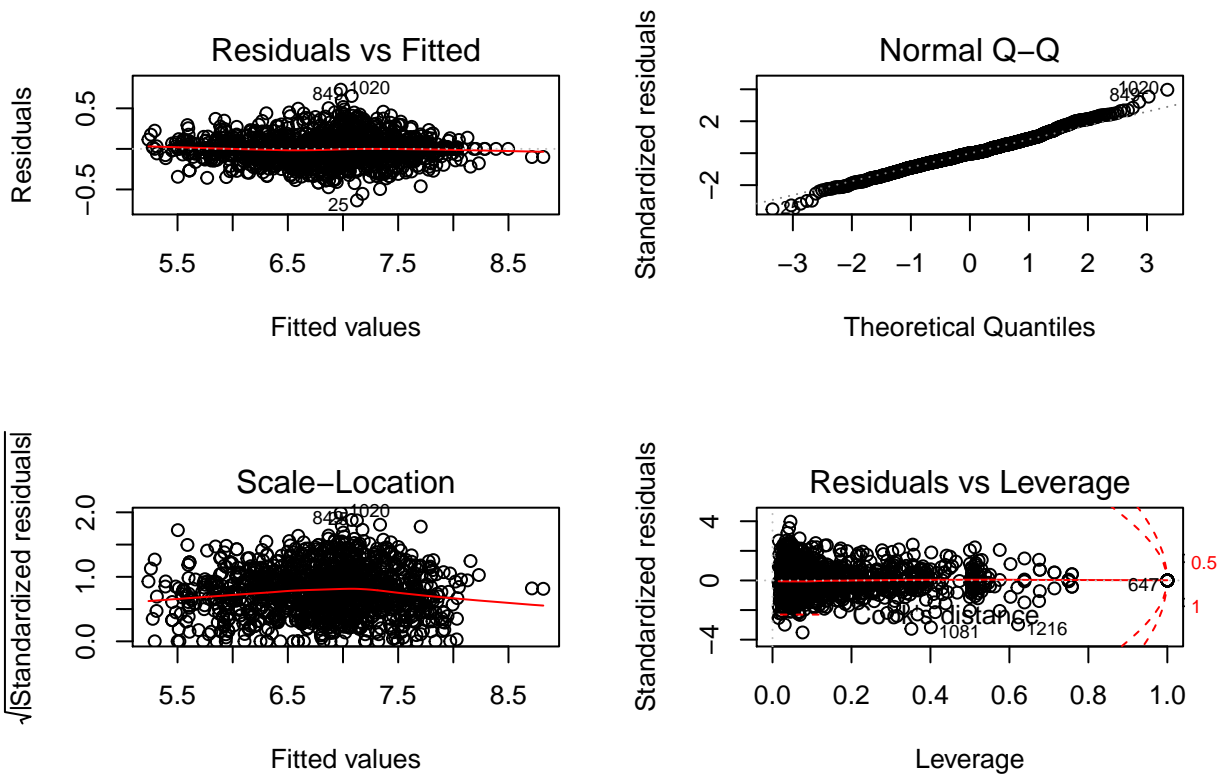
```
## 13, 15, 18, 29, 34, 46, 84, 128, 160, 173, 178, 179, 205, 232, 267, 271, 299, 303, 324, 348, 388, 4
```

```
## Warning: not plotting observations with leverage one:
```

```
## 13, 15, 18, 29, 34, 46, 84, 128, 160, 173, 178, 179, 205, 232, 267, 271, 299, 303, 324, 348, 388, 4
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```

```
ad.test(lm_full_t$residuals) #not really
```

```
##
## Anderson-Darling normality test
##
## data:  lm_full_t$residuals
## A = 7.5169, p-value < 2.2e-16
```

```
shapiro.test(lm_full_t$residuals) #not really
```


```
##
## Shapiro-Wilk normality test
##
## data:  lm_full_t$residuals
## W = 0.98478, p-value = 1.874e-10
```

A look over outliers

```
library(car)
```

```
## Loading required package: carData
##
## Attaching package: 'car'
## The following object is masked from 'package:EnvStats':
##
##     qqPlot
## The following object is masked from 'package:psych':
##
##     logit
```

```
influencePlot(lm_full,main="Influence Plot", sub="Circle size is propoertial to Cook's Distance" )
```



```
##      StudRes      Hat      CookD
## 13      NaN 1.00000000      NaN
## 15      NaN 1.00000000      NaN
## 434  2.124383 0.67519240 0.030955546
## 849  4.607995 0.04198279 0.003020118
## 1020 5.093056 0.04545388 0.003990584
## 1081 -4.381193 0.40114343 0.041814955
```

```
#Cook's Distance
```


```
cooksds <- cooks.distance(lm_full_t)
cooksda=data.frame(cooksds)
summary(cooksds)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
## 0.00000 0.00003 0.00013 0.00076 0.00053 0.04815      88
```


```
# identify D values > 4/(n-k-1)
```

```
# Cook's D plot
```

```
cutoff <- 4/((nrow(data)-length(lm_full_t$coefficients)-2))
plot(lm_full_t, which=4, cook.levels=cutoff)
```



```
plot(cooksd, pch="*", cex=1, main="Influential Obs by Cooks distance") # plot cook's distance
abline(h = cutoff, col="red") # add cutoff line
text(x=1:length(cooksd)+1, y=cooksd, labels=ifelse(cooksd>4*mean(cooksd, na.rm=T), names(cooksd), ""),
col="red") # add labels
```



```

#extract influential obs
influential <- as.numeric(names(cooksdata)[(cooksdata > cutoff)]) # influential row numbers
influ=data.frame(data[cooksdata > cutoff, ])
filtered_data <- data[ !(row.names(data) %in% c(influential)), ]
#Outlier rimossi
lm_full_t_no_OUTliers = lm(log(Price) ~ ., data = filtered_data)
par(mfrow=c(2,2))
plot(lm_full_t_no_OUTliers)

```

```
## Warning: not plotting observations with leverage one:
```

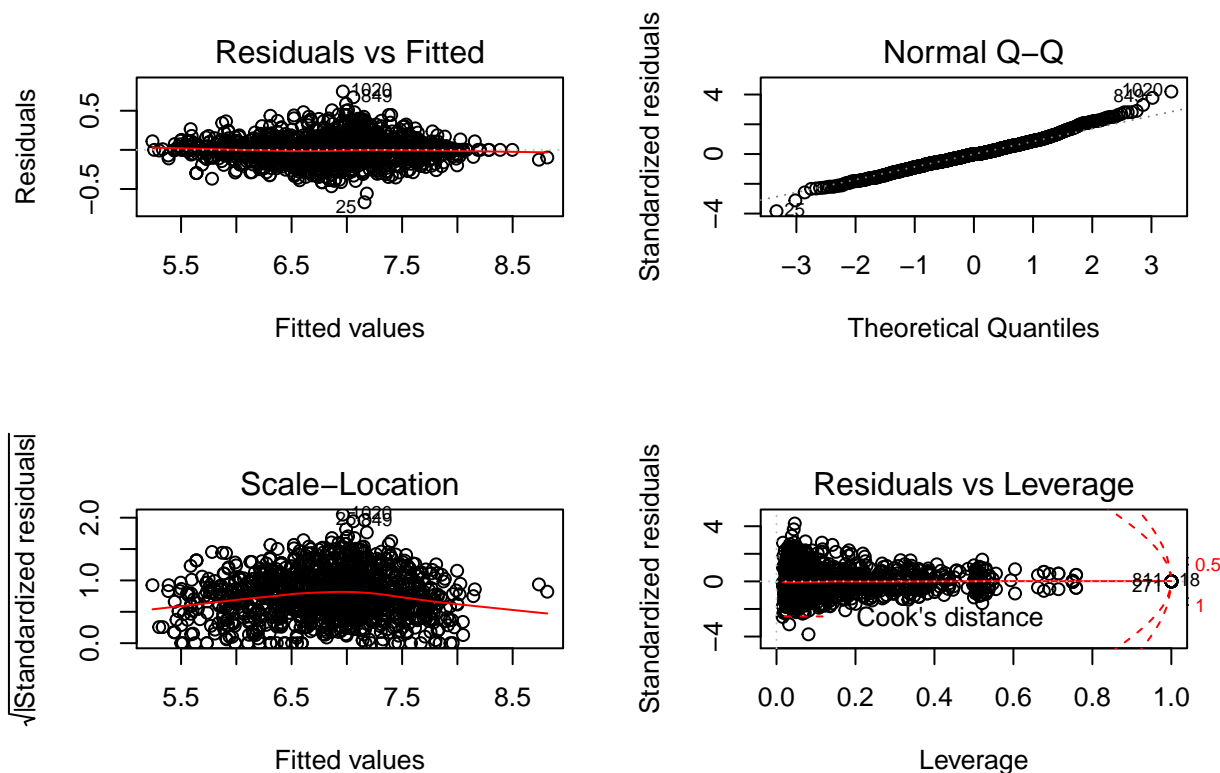
```
## 15, 21, 32, 43, 148, 172, 196, 222, 288, 292, 312, 335, 408, 420, 430, 438, 439, 447, 456, 481, 48
```

```
## Warning: not plotting observations with leverage one:
```

```
## 15, 21, 32, 43, 148, 172, 196, 222, 288, 292, 312, 335, 408, 420, 430, 438, 439, 447, 456, 481, 48
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```

```
## Warning in sqrt(crit * p * (1 - hh)/hh): NaNs produced
```



```

#summary(lm_full_t_no_OUTliers) #FIXME: too long to be printed, R^2=0.9727
ncvTest(lm_full_t_no_OUTliers)

```

```
## Non-constant Variance Score Test
```

```
## Variance formula: ~ fitted.values
```

```
## Chisquare = 1.740444, Df = 1, p = 0.18708
```

```
null = lm(log(Price) ~ 1, data = filtered_data)
```

```
full = lm(log(Price) ~ ., data = filtered_data)
```

```
lm_fit = stepAIC(null, scope = list(upper = full), direction = "both", trace = FALSE)
```

```
drop1(lm_fit, test = 'F')
```

```
## Single term deletions
##
## Model:
## log(Price) ~ Cpu + Memory + OpSys + Gpu + TypeName + ScreenResolution +
##      Company + Ram + Inches
##
```

	Df	Sum of Sq	RSS	AIC	F value	Pr(>F)
## <none>			32.161	-4053.8		
## Cpu	84	15.5056	47.667	-3724.9	5.5789	< 2.2e-16 ***
## Memory	34	9.9754	42.136	-3780.6	8.8673	< 2.2e-16 ***
## OpSys	5	5.9181	38.079	-3850.5	35.7726	< 2.2e-16 ***
## Gpu	85	10.7290	42.890	-3860.2	3.8149	< 2.2e-16 ***
## TypeName	4	2.1668	34.328	-3979.5	16.3717	5.483e-13 ***
## ScreenResolution	28	4.7325	36.893	-3936.4	5.1082	5.348e-16 ***
## Company	14	3.0807	35.242	-3966.3	6.6506	4.801e-13 ***
## Ram	1	1.5499	33.711	-3996.4	46.8425	1.360e-11 ***
## Inches	1	1.0393	33.200	-4015.7	31.4100	2.720e-08 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```